

Take-Home Task — Gong → Salesforce Draft Agent

Candidate: Luka Cerrutti

Role: AI Builder, Revenue AI team @ Tastewise

One-line summary

Build a tool-using agent that reads a mock sales call transcript and produces a structured Salesforce opportunity update draft, with explainable reasoning, typed output, and an explicit abstain path for when the transcript doesn't contain enough signal.

Why this matters

Tastewise AEs spend ~30 minutes per call cleaning up Salesforce afterward — MEDDPIC fields, next steps, stage changes, close-date shifts. Our Revenue AI team is building the agent that collapses this to a 30-second human review.

Your call stood out — LangChain + Semantic Kernel + Claude Agent SDK seminars at EY, plus a maintained side product (nexata) with real paying clients (hospital + Huracán C-level). This task lets you apply the same multi-agent pattern you built at EY's `superaciones` platform (the NL → JSON → Excel → validation → variance-report flow you walked me through at [21:06]) to a GTM use case. Stretch for you: framing the system around a revenue workflow rather than an accounting one.

What to build

An **agent** (not just an LLM call) that:

1. **Ingests** a mock Gong transcript (plain text, provided below).
2. **Extracts** structured signal:
 - Deal stage delta (e.g. *"Discovery → Solution Design"*).
 - Next step + owner + date.

- Champion / economic buyer / detractors mentioned.

- Risks + objections raised.

- Suggested close-date shift (with reasoning).

3. **Drafts** a Salesforce opportunity update as a structured JSON payload matching the Opportunity schema below.

4. **Explains** its reasoning — for each updated field, cite the transcript quote + timestamp that justifies the change.

5. **Knows when to abstain** — if there's no real signal (mostly small talk), it says so and proposes nothing, while still logging a last-touch summary.

Required agentic properties

- It must use **tools** — not just one LLM prompt. At minimum: a *transcript-reader* tool and a *schema-validator* tool. Bonus: an *uncertainty-check* tool that flags low-confidence fields.
- It must handle **both test transcripts** (rich + thin — provided below).
- It must produce **typed output** that would pass basic schema validation against the SFDC-lite schema.

Stack

Pick any agent framework — **Claude Agent SDK, LangGraph, LangChain, AutoGen, PydanticAI, CrewAI, Semantic Kernel, hand-rolled with structured outputs**. Given what you described at EY, you have real depth in at least three of these — we want to see you pick deliberately and explain the trade-off. Extra credit for naming a rejected alternative and why.

Do not use a real Salesforce instance or real Gong data. Mocks only.

Deliverables

1. **Source code** (GitHub repo or zip) — minimal is fine. We want to see the agent loop, tool definitions, and the schema.

2. **A short Loom or screen recording** (≤5 min) running the agent on both test transcripts, showing:

- The reasoning trace (or close to it).
- The final JSON output for Transcript 1.
- The abstain behavior on Transcript 2.

3. **A write-up** (≤ 1.5 pages). Cover:

- Framework choice and why — what did you reject?
- How you prompt/structure the extractor vs the drafter — are they one call or two?
- Which model(s) you picked for each step and why.
- How you'd evaluate this in production — your eval harness concept.
- Failure modes you'd expect in the wild + mitigations.
- What you cut and would add with another 4 hours.
- **Bonus:** one thing you learned at EY about multi-agent orchestration that you'd apply (or explicitly *not* apply) here.

Explicit scope guardrails

- **No UI needed.** CLI or notebook is fine.
- **No deployment.** Run local.
- **No fine-tuning, no RAG, no vector DB** — wrong scope for this task.
- **Cut scope before you run over.** A clean 80% demo beats a messy 100%.

How we'll evaluate

Not a secret — we're grading:

1. ****Does the agent actually *decide* things? Tool-calling, branching, abstaining — not a single chained prompt.**
2. **Is the output trustworthy? — Typed, grounded in quotes, honest about uncertainty.**
3. **Is the schema handling real? — Would the JSON pass a real SFDC validator?**
4. **Production reasoning in the write-up — Can you operate this, not just build it?**
5. **Shipping discipline — Working end-to-end > sprawling + broken.**

Provided: mock SFDC-lite Opportunity schema

```
``json
```

```
{
```

```
"opportunity_id": "string (required)",
```

```

"stage": "enum: prospecting | discovery | solution_design | proposal | negotiation |
closed_won | closed_lost",

"amount_usd": "number",

"close_date": "ISO date string",

"next_step": {"description": "string", "owner": "string", "due_date": "ISO date
string"},

"meddpicc": {

"metrics": "string", "economic_buyer": "string", "decision_criteria": "string",

"decision_process": "string", "paper_process": "string", "identify_pain": "string",

"champion": "string", "competition": "string"

},

"risks": ["string"],

"last_touch_summary": "string",

"confidence": "float 0-1"

}
,

```

Provided: mock transcripts

Transcript 1 – "Rich signal" (~40 lines, condensed)

```

,

[00:02] AE: Thanks for the time, Priya. Last time we left off you were going to loop
in Ben from Procurement.

[00:18] Priya: Yeah – Ben's in the next call. He's fine with the $120k range but wants
annual not quarterly.

[00:34] AE: Understood. And on the evaluation criteria – is speed-to-insight still #1?

[00:45] Priya: Speed is #1. Second is how well it integrates with our innovation
stage-gate. We'd want that in writing.

[01:22] AE: Noted. On timing – still targeting Q3 kickoff?

[01:30] Priya: Actually pushing to early Q4 now. Our 2027 portfolio review moved.

[02:05] AE: Okay. So on my side – I'll send a revised quote tonight annual, and a
one-pager on stage-gate integration by Friday. You'll loop in Ben for next Tuesday?

```

[02:20] Priya: Yes. Ben, me, and our CMO. CMO will be the signer.

[02:45] AE: Perfect. Any concerns from your side?

[02:55] Priya: One. I saw Spoonshot in market last week. The demo was lightweight but the price was aggressive. You'll need to be ready for that.

,

Transcript 2 – "Thin signal" (~15 lines)

,

[00:02] AE: Hey! How was the trip?

[00:06] Champion: Oh amazing, went to Lisbon. You been?

[00:10] AE: Never. On the list.

[00:14] Champion: Highly recommend. Anyway – I know we had this on the calendar, but honestly I don't have much new. Still waiting to hear back from finance.

[00:40] AE: Totally fine. Any sense of timing?

[00:44] Champion: Maybe two weeks? I'll ping you.

[00:50] AE: Sounds good. Keep me posted.

..

The agent should produce a meaningful update for #1 and propose no update for #2 (while still logging the touch).

Questions?

Email Nico at nicolas@tastewise.io. If a question would change the shape of what you build, ask – don't guess silently.

The nexata walkthrough stuck with me. The fact that you maintain a shipped product with paying clients alongside EY is a rare signal – I want this task to show me you can re-apply that pattern to the GTM domain we're building for.

– Nico